

# Information Retrieval

2022/2023

## Assignment 1

Submission deadline: **21 October 2022**

For this assignment you will implement a text indexer following the SPIMI approach.

The corpus for this assignment was compiled from the ‘BioASQ Task B Dataset’, described here <http://bioasq.org/participate/challenges>

1. Create a corpus reader that iterates over the documents and returns, in turn, the contents of each document.  
For this assignment consider only the `title` and `abstract` as content (indexed) fields and use the `pmid` as the document identifier.
2. Create a tokenizer based on tokenizing rules specified by you. It should also include different settings to allow adjusting its behaviour. As minimal requirements, your tokenizer should allow the user to apply:
  - i. A minimum length filter that removes tokens with less than the specified number of characters;
  - ii. Normalization to lowercase;
  - iii. A stopwords filter, allowing the user to specify a file listing the stopwords to use;
  - iv. Porter stemmer (from Snowball, using the `PyStemmer` wrapper, or from NLTK).
3. Implement an indexer that follows the SPIMI approach. To signal the dumping of indexing blocks to disk you may consider either: 1) the amount (or %) of available memory; 2) the memory required by the index, estimated by the number of posting and the size of the posting in bytes.  
Note: You should experiment with different values for strategy 1) or 2), to simulate a system with lower memory resources.
4. Index, separately, each of the files listed below and gather the following statistics:
  - a) Total indexing time
  - b) Merging time (last SPIMI step)
  - c) Number of temporary index segments written to disk (before merging)
  - d) Total index size on disk
  - e) Vocabulary size (number of terms)

Files to use:

pubmed\_2022\_tiny.jsonl.gz (128 MB)  
pubmed\_2022\_small.jsonl.gz (1.3 GB)  
pubmed\_2022\_medium.jsonl.gz (4.1 GB)  
pubmed\_2022\_large.jsonl.gz (8.8 GB)

**Instructions:**

- **Modelling**, code **structure**, **organization** and **readability** will be considered when grading your project
- **Comment** your code; and make sure you include your name and student number
- Write **modular** code
- Favour **efficient** data structures
- Use **parameters**, preferably through the command line
- Make sure all your programs run correctly
- Submit your assignment by the due date using Moodle